# Lab 2: Fine-tune vs prompt

## Does fine-tuning beat your best prompt?

Aarne Talman

5 November 2026

LDA-T318 | LLMs and Generative AI for NLP | University of Helsinki

# The question

## Lab 2

For *your* task, does fine-tuning beat your best prompt?

You compare two ways of getting a model to do your task:

- **prompting**: give the model instructions and, optionally, a few examples, without changing it
- **fine-tuning**: train the model on your labelled examples

Both are scored on the same test examples, which were not used for anything else.

### For the final report

- one results table (`summarize_results.py` prints it)
- one loss plot
- one paragraph: was fine-tuning worth the extra time and GPU use?

**Minimum:** README steps 1–11 once, with default settings. Ideas for going further are in the README.

Everything you need is in `Week-2/Lab-2/README.md`, step by step.

# Today

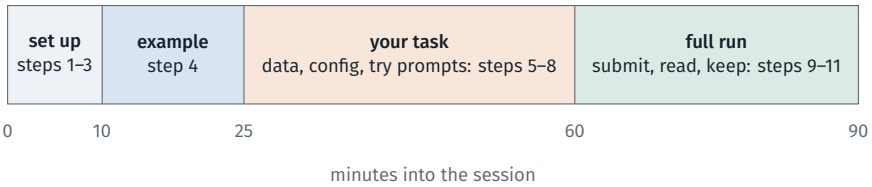


The steps are those in the README. Start with the **smoke test** (step 2): it checks the scripts without a GPU. If it has not passed by minute 15, put your hand up.

The full experiment is **one GPU job**. It scores the prompts, trains the model and runs the general-knowledge check, and it can take an hour or more. Submit it by minute 60, even with a first version of your prompt.

GPUs are reserved until about 20:00, so jobs can finish after the session.

# Your data and your prompt

```
# 1. download an existing labelled dataset and split it (README step 5)
python3 from_huggingface.py --dataset stanfordnlp/sst2 --out sst2.jsonl
python3 make_splits.py --input sst2.jsonl --fields sentence \
    --label-field label --out data/my-task

# 2. lab2lib/prompting.py: the two settings you change (step 6)
TASK_FIELDS = ("sentence",)
TASK_SYSTEM = ("Is the review positive or negative? "
               "Use exactly one of: positive, negative. "
               "End with 'Label: <label>'.")
```

**Three splits of your data**

- **train**: what the model learns from (1,000–5,000)
- **dev**: for choosing the prompt and settings
- **test**: only for the final scores

You do **not** create a dataset. Use an existing labelled one: SST-2 (above), another dataset listed in the README, or your own file.

Keep the last sentence of the instruction: the scorer looks for `Label: <label>`. In step 8 you try the instruction on dev examples with `try_prompt.py` and improve it.

# The experiment

## One GPU job does everything

You prepare the data and config (README steps 5–7) and try your prompt interactively on dev examples (step 8). Then you submit **one GPU job** (step 9), which runs the five phases below in order, and read the results (step 10).

| | |
|---|---|
| **1. Prompting** | your instruction in four versions: base and instruct model, each with no examples and with 5 examples; the best on dev is *the* baseline |
| **2. Knowledge, before** | the instruct model answers 50 general-knowledge questions |
| **3. Fine-tuning** | LoRA trains a small add-on to the instruct model on your training data |
| **4. Testing** | the best prompt and the fine-tuned model are scored on the same test set |
| **5. Knowledge, after** | the fine-tuned model answers the same 50 questions |

**Your conclusion** (step 11): does fine-tuning beat the best prompt, and at what cost?

**Read the outputs before you trust a score:** a score of exactly 0 or 100 usually means the scorer is wrong, not the model (`show_examples.py`).

# If you have time: tune the training

| Setting | Start | What it controls |
|---|---|---|
| learning rate | 0.0002 | how large each training step is |
| epochs | 2 | how many times the model sees the training data |
| LoRA rank | 16 | the size of the adapter |

Change *one* setting at a time, give each run its own name, and keep a small table of the runs. The README lists other extensions.

**Lower training loss is not a result**

The **loss** measures how wrong the model's answers are. If the training loss falls while the dev loss rises, the model is **memorising** the training examples.

Then try fewer epochs or a lower learning rate. Choose settings on the dev loss, never on the test score.

# Checklist for today

[ ] smoke test passed (step 2)
[ ] example job finished on `gputest` (step 4)
[ ] dataset downloaded and split (step 5)
[ ] prompt adapted (step 6)
[ ] config made, `check_data.py` clean (step 7)
[ ] prompt tried on dev with `try_prompt.py` (step 8)
[ ] full job submitted to `gpumedium` (step 9)
[ ] table printed with `summarize_results.py`
[ ] ten raw outputs read with `show_examples.py`
[ ] results folder and config kept for the report (step 11)

## Good practice on Roihu

- follow a job with `squeue --me` and `tail -f logs/lab2-lora-<jobid>.out`
- do not run models on the login node
- leave interactive GPU sessions with `exit` when you are done, and close unused GPU Jupyter sessions

Stuck for more than ten minutes? Put your hand up. The README has a troubleshooting table.